



DATABASE INDEXES



#SperasoftTalks

Sperasoft is a great team of professionals specializing in SKU development, game engineering, art, online and platform development and production support.

***follow us on twitter
@Sperasoft***


visit our site:
<http://sperasoft.com/career/>




- RDBMS store data **only** in Trees
- Index is a tree in terms of data structure
- a **Table** is an **Index**
- a **Clustered Index** is a **Table** itself
- a Non-clustered Index is a **copy** of data
- all Non-clustered Indexes refer to **Clustered** one
- all keys in Tree Nodes are always **unique**

What's common between

- 
- Oracle Database
 - SQL Server
 - IBM DB2
 - MySQL
 - PostgreSQL
 - Sybase
 - Informix

- 
- RDBMS is a type of Database Management System that stores data in the form of related tables
 - RDBMS is a Database Management System that is based on the relational model introduced by E.F. Codd
 - Data is stored in tables and the relationships among the data are also stored in tables

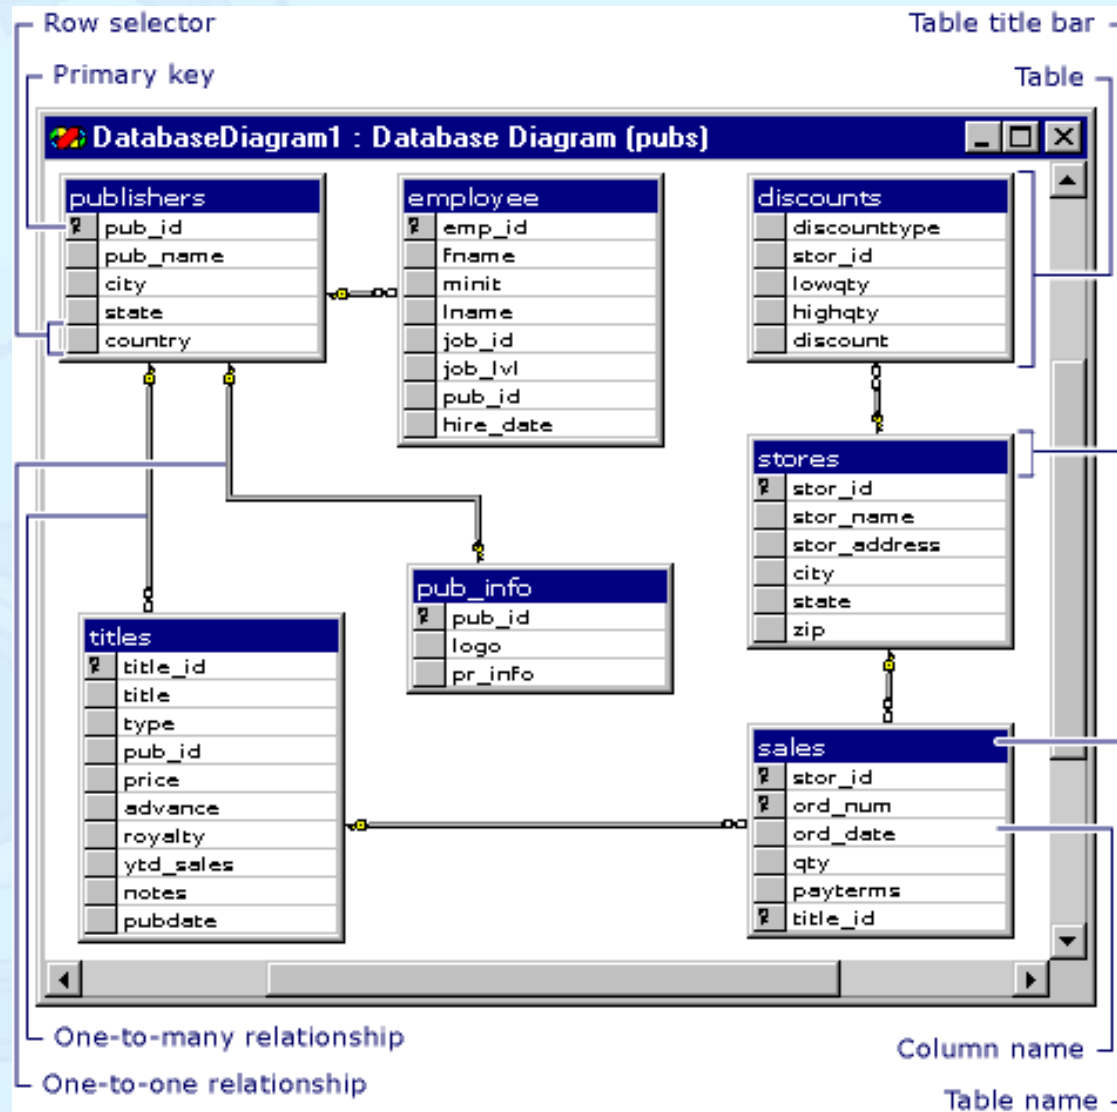


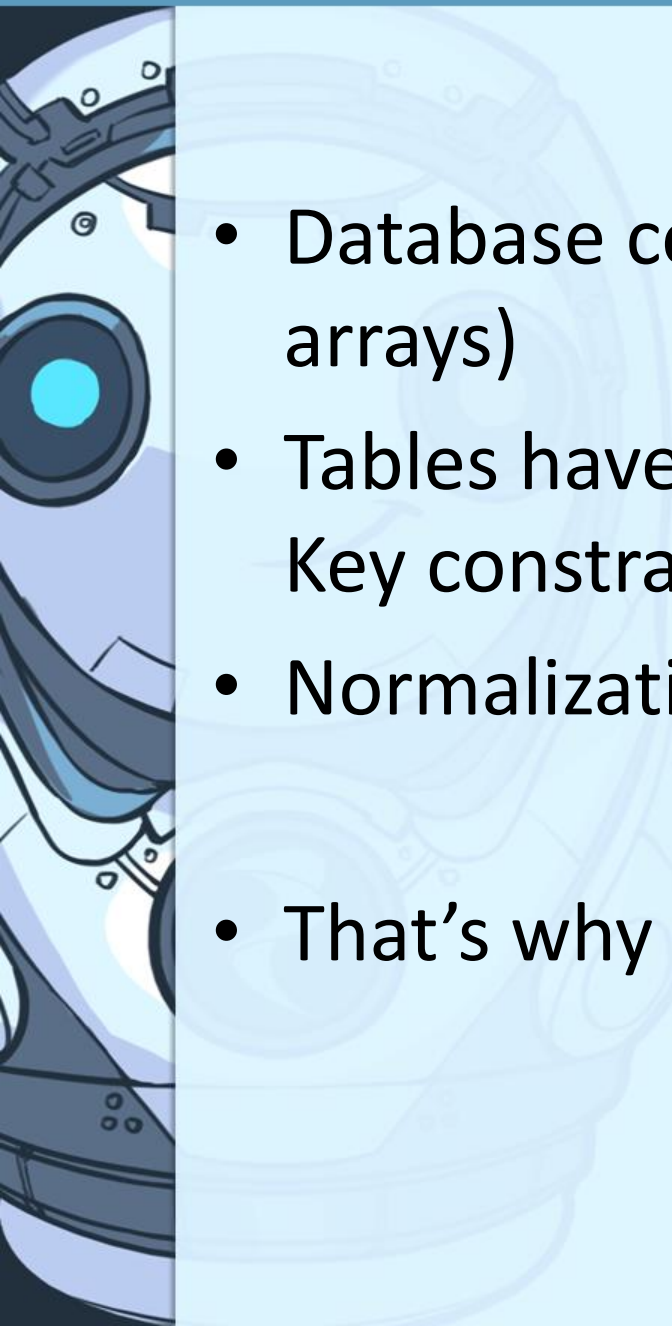
- Born on the Isle of Portland in England in 1923
- Died in Florida US in 2003, aged 79
- Mathematic
- Worked for IBM

- Introduced “A Relational Model of Data for Large Shared Data Banks” and Alpha database language
- IBM started implementing the Relational model and introduced another language named SEQUEL

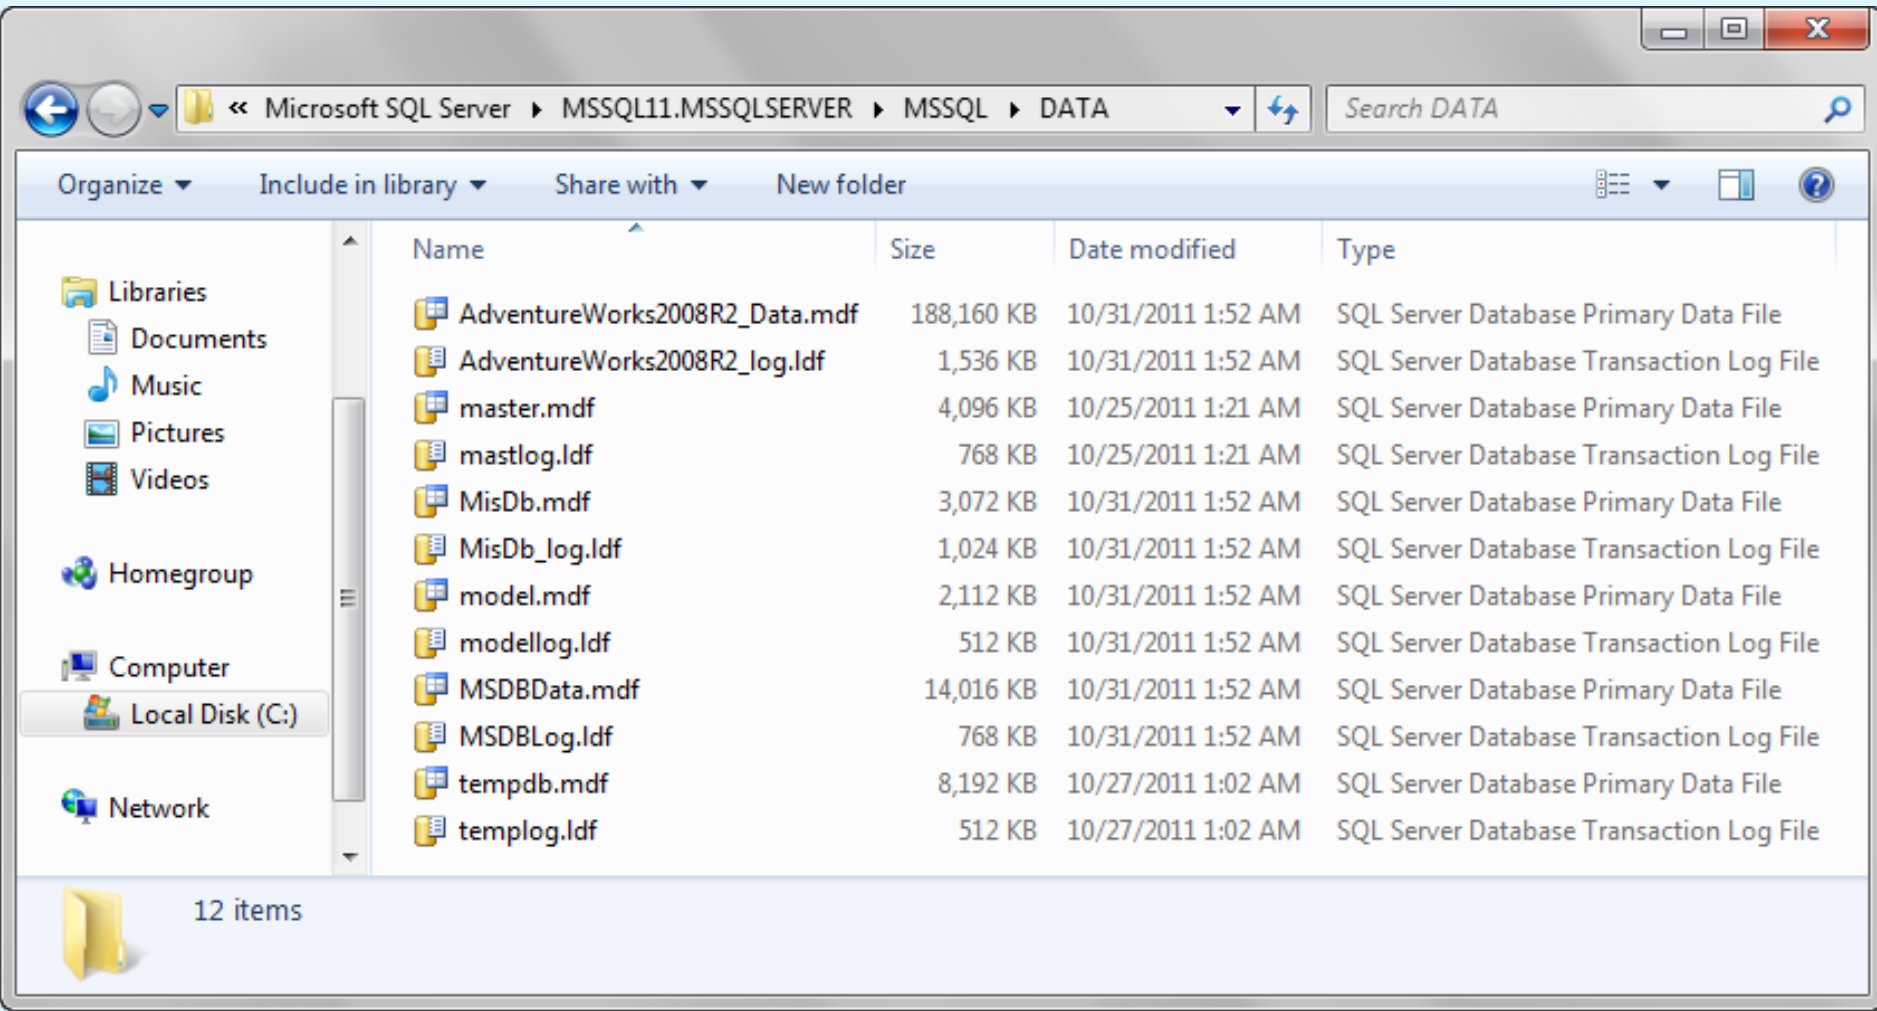
- Larry Ellison came up in time with his implementation of Relational model and the language – Oracle Database and SQL
- ANSI started making SQL standard

- It's all about Table Relations




- 
- The slide features a light blue background with a faint, stylized illustration of a robot's head in the center. On the left side, there is a vertical strip showing a close-up of a robot's eye and mechanical components. The text is presented in a clean, black, sans-serif font.
- Database contains tables (two dimensional arrays)
 - Tables have relationships enforced by Foreign Key constraints (1-to-Many relationship)
 - Normalization of tables is a key concept
 - That's why RDBMS are called Relational


What's database physically



All tables are stored in a file

- Files are flat in nature

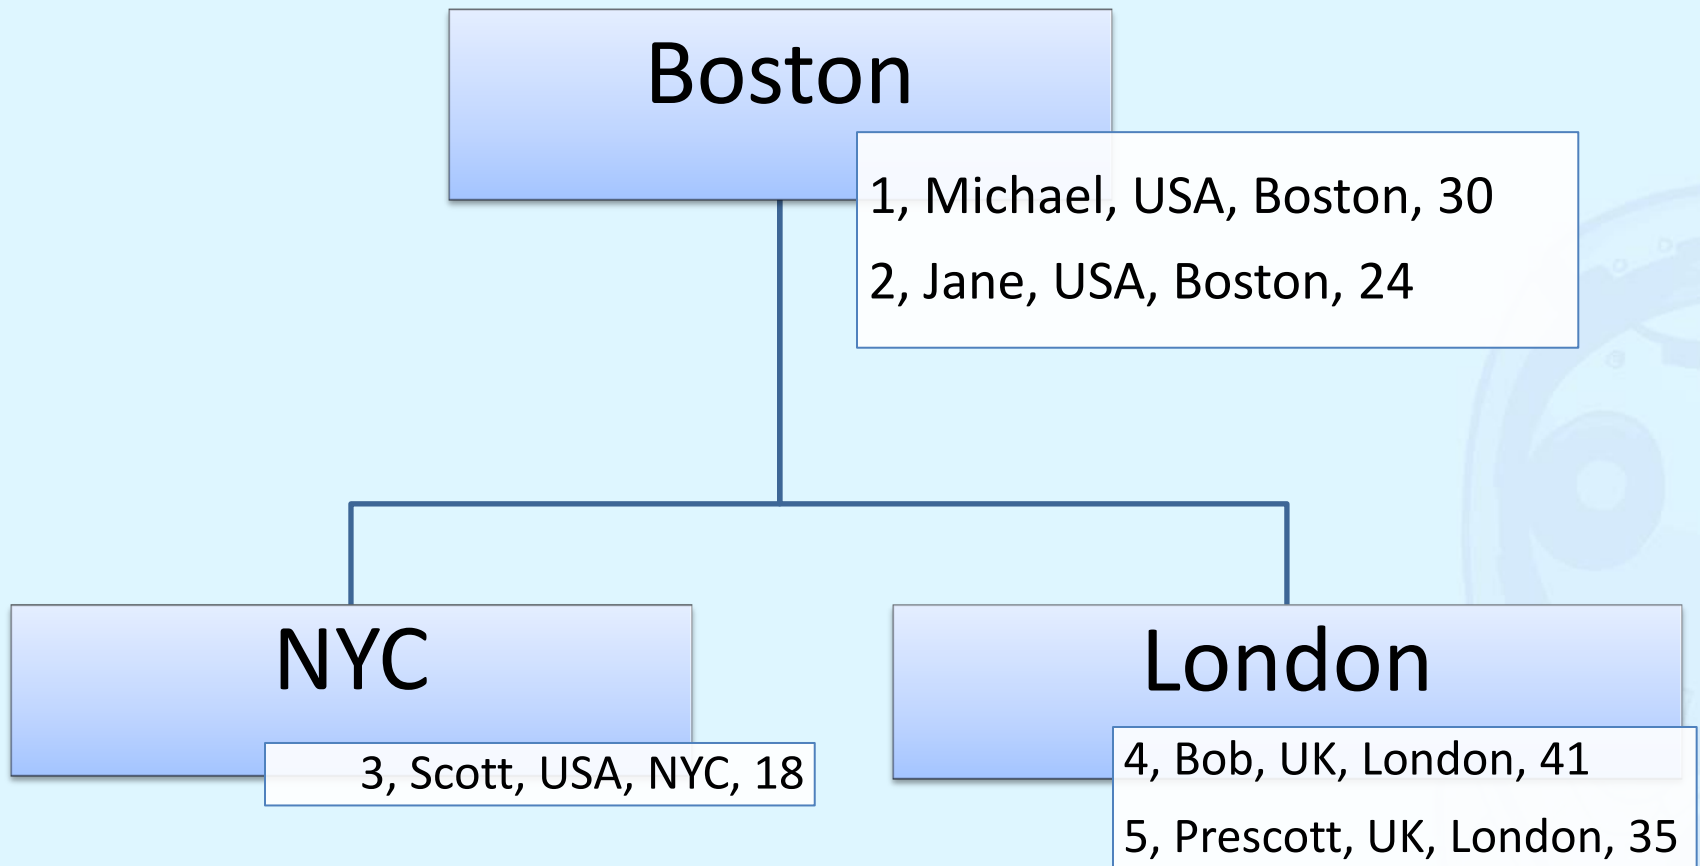


- 
- What's the value behind relations?
 - What is a database table?
 - What is a table index?
 - Relations vs How data is stored


```
ArrayList<User> users = new ArrayList<User>();
```

- Such array seems to be a table
- How to find Users from Boston faster?

Id	User Name	Country	City	Age
1	Michael	USA	Boston	30
2	Jane	USA	Boston	24
3	Scott	USA	NYC	18
4	Bob	UK	London	41
5	Prescott	UK	London	35

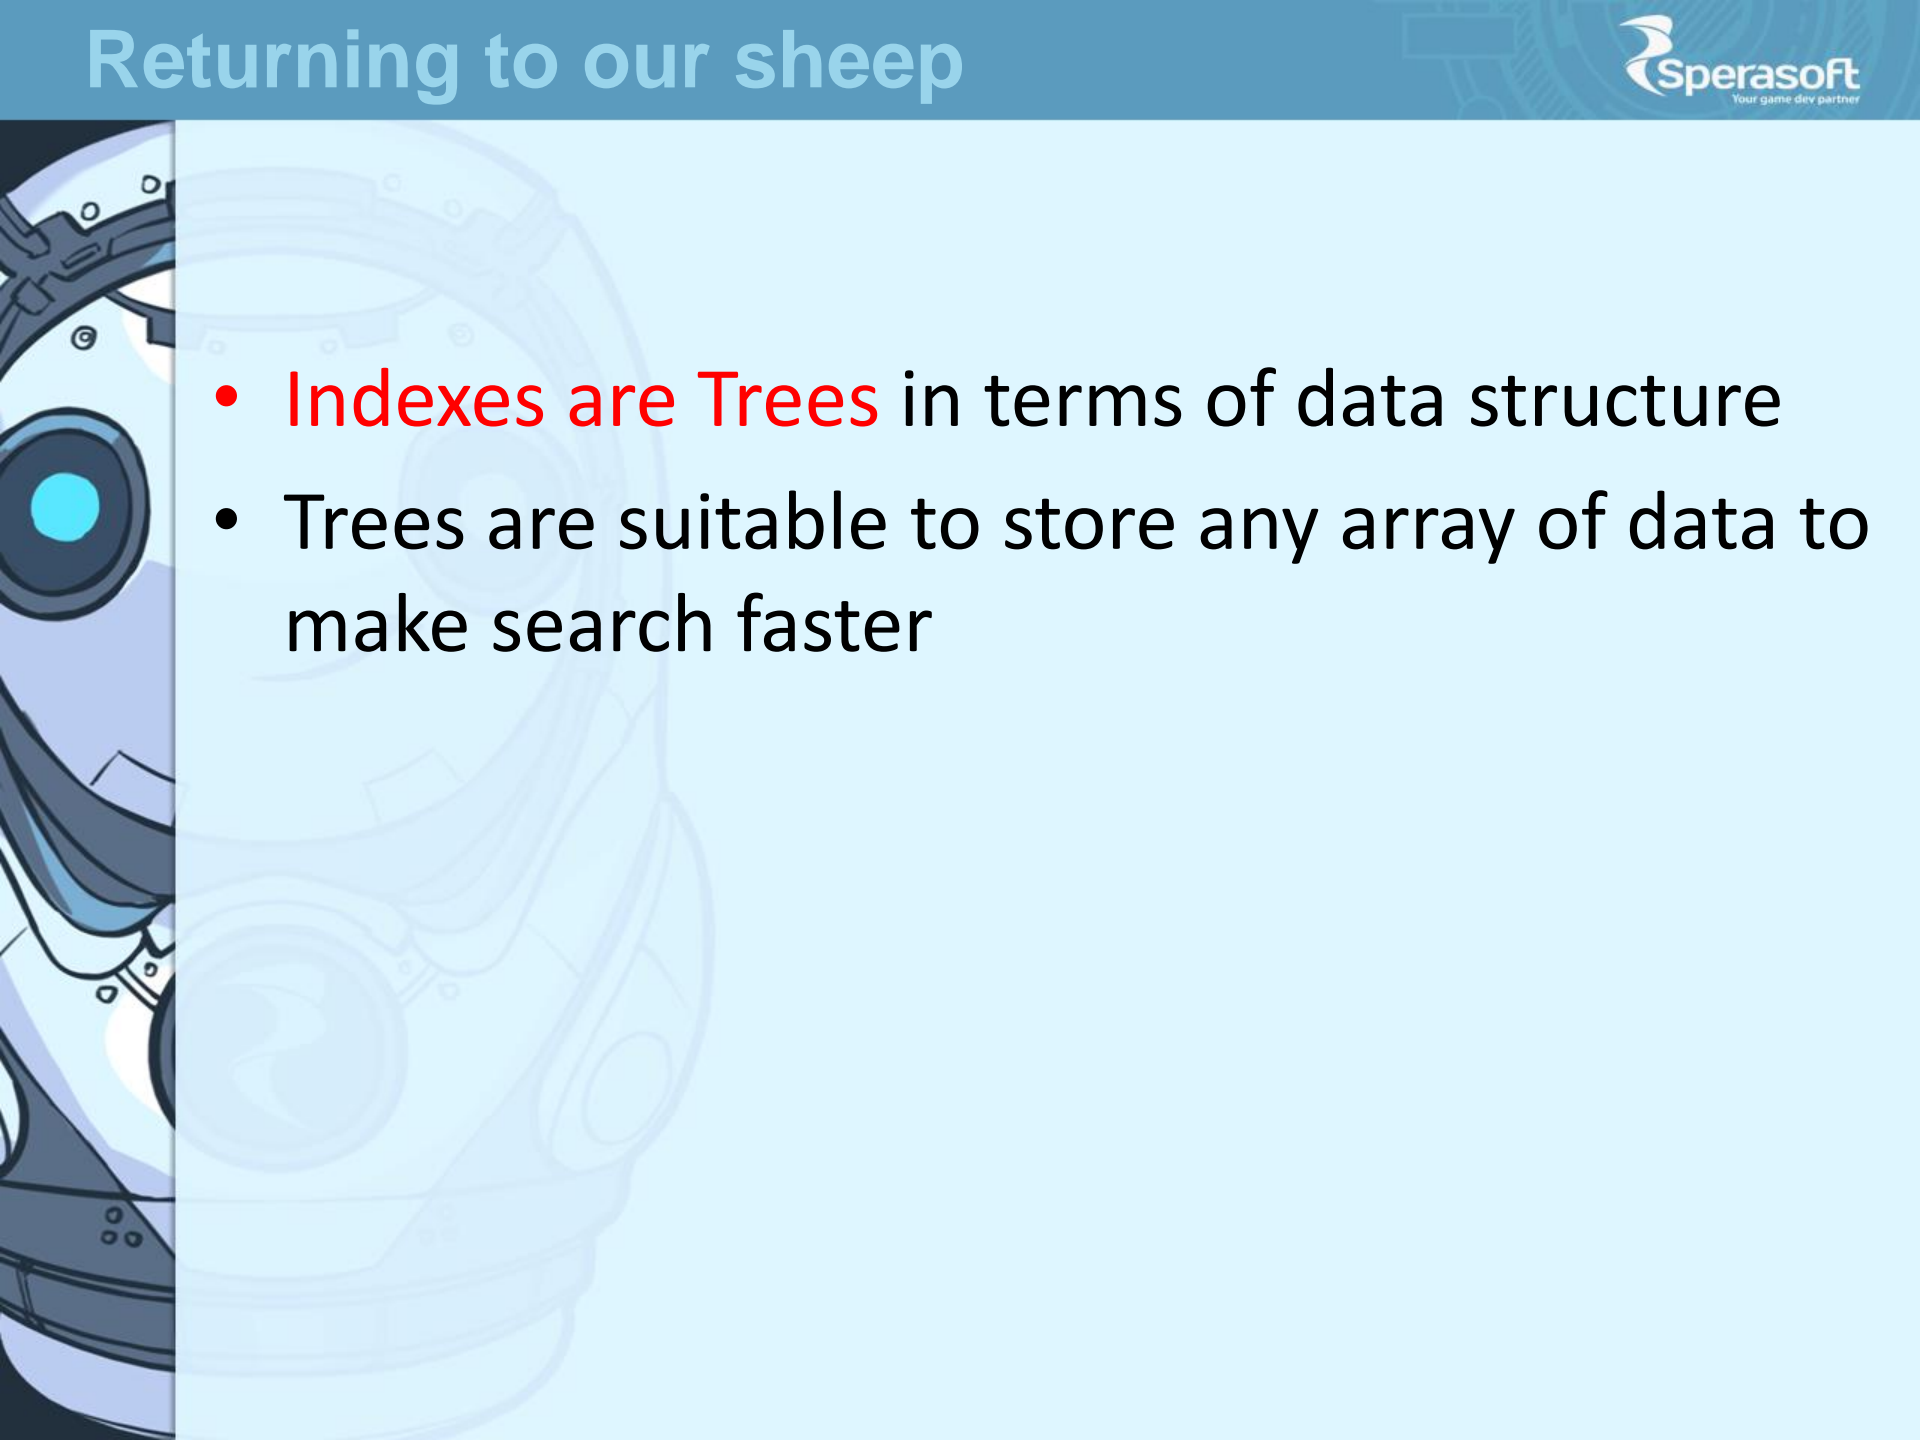


Can replace an initial array with Index

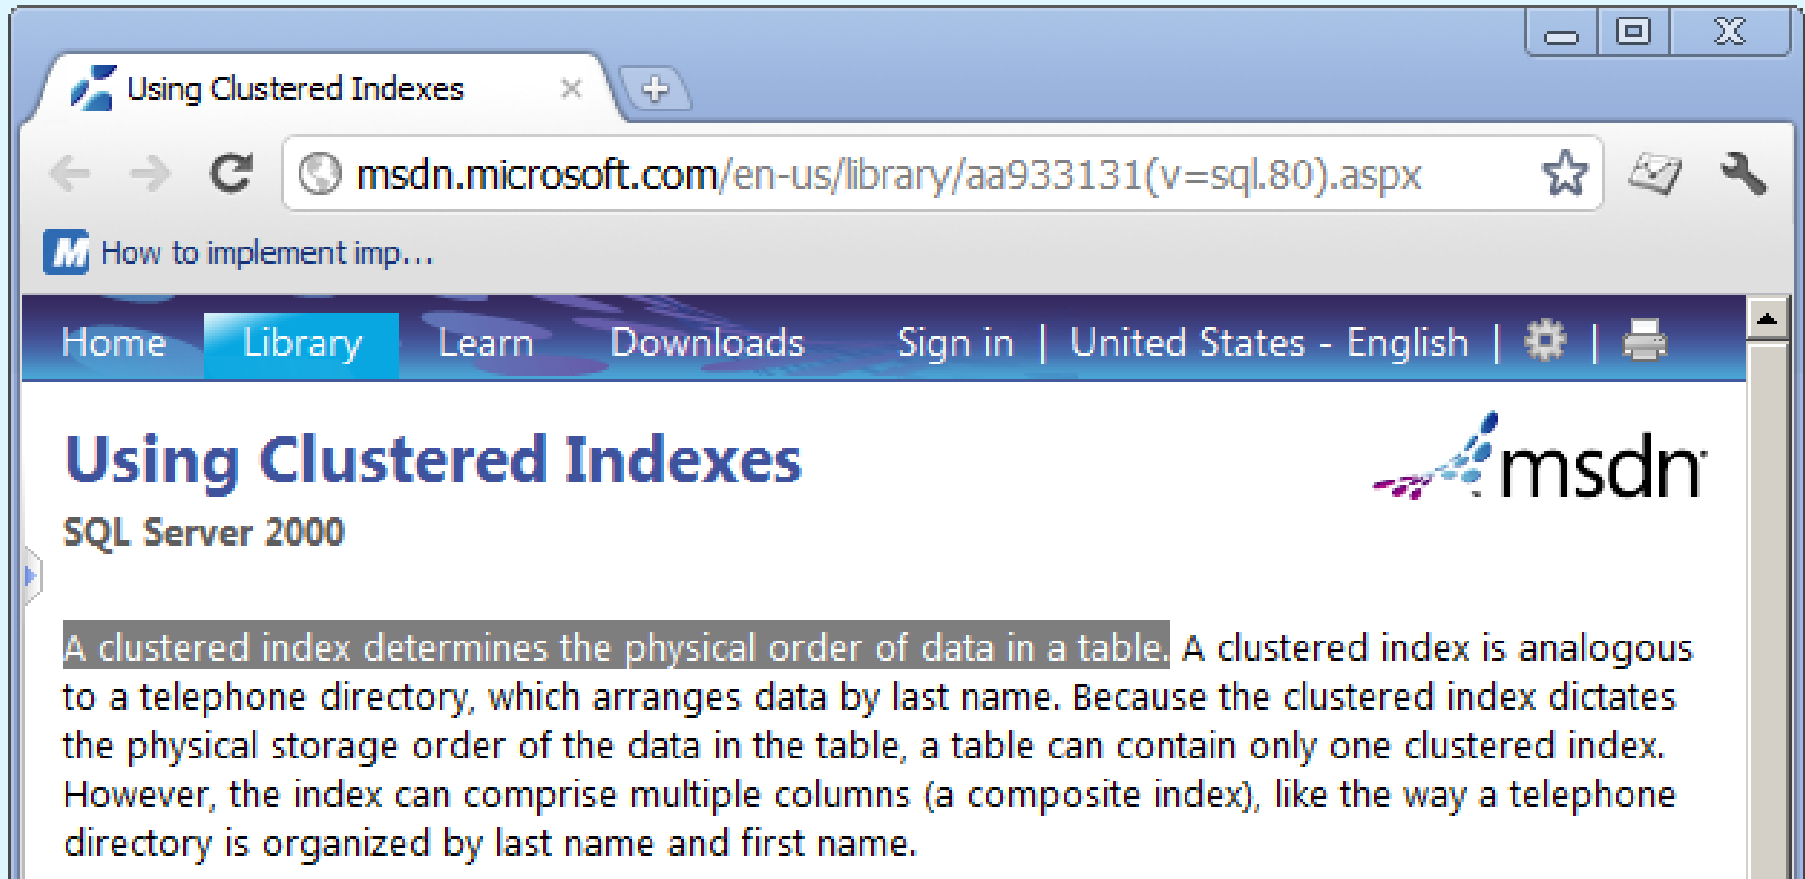


What's important to note

- Key values in a Key node should be **unique**
- Otherwise Trees do not work

- 
- **Indexes are Trees** in terms of data structure
 - Trees are suitable to store any array of data to make search faster

- All RDBMS store data as **Balanced Trees**
- The concrete implementation of B-Tree could differ from vendor to vendor
- It means the only way to store data is **Tree**
- No exceptions here - **table is a tree, index is a tree**



The screenshot shows a web browser window with the title 'Using Clustered Indexes'. The address bar displays the URL `msdn.microsoft.com/en-us/library/aa933131(v=sql.80).aspx`. The browser's address bar shows the URL `msdn.microsoft.com/en-us/library/aa933131(v=sql.80).aspx`. The page header includes navigation links: Home, Library (selected), Learn, Downloads, Sign in, United States - English, and icons for settings and printing. The main content area features the title 'Using Clustered Indexes' and the subtitle 'SQL Server 2000'. The MSDN logo is visible in the top right corner of the content area. The main text explains that a clustered index determines the physical order of data in a table, is analogous to a telephone directory, and that a table can contain only one clustered index. It also mentions that the index can be a composite index.

Using Clustered Indexes

SQL Server 2000

A clustered index determines the physical order of data in a table. A clustered index is analogous to a telephone directory, which arranges data by last name. Because the clustered index dictates the physical storage order of the data in the table, a table can contain only one clustered index. However, the index can comprise multiple columns (a composite index), like the way a telephone directory is organized by last name and first name.

- The next record in Clustered Index is always stored after the previous one

RECORD 1	RECORD 2
1 Michael USA Boston 30	2 Jane USA Boston 24



**Have a question?
Like this deck?**

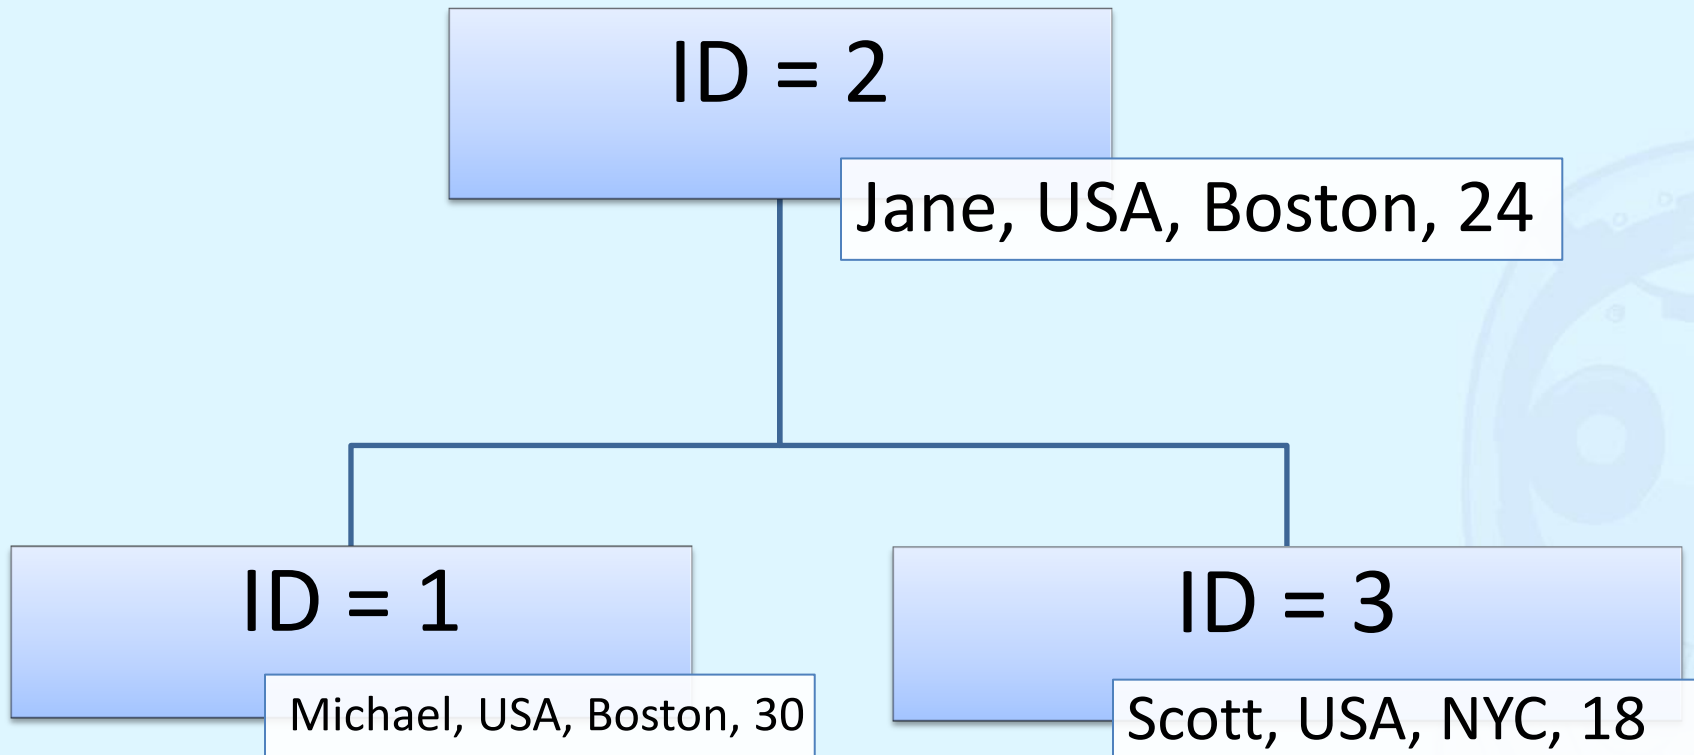
Just follow us on twitter @Sperasoft

- Clustered Indexes
 - Non-clustered indexes
 - Both could be unique and non-unique
 - Table can be without any indexes
-
- How is that comply with how data is actually stored?

- Unique and non-unique
- CREATE CLUSTERED INDEX [name] ON [table_name] ([column1], [column2])
- CREATE **UNIQUE** CLUSTERED INDEX [name] ON [table_name] ([column1], [column2])

- Unique and non-unique
- CREATE NONCLUSTERED INDEX [name] ON [table_name] ([column1], [column2])
- CREATE **UNIQUE** NONCLUSTERED INDEX [name] ON [table_name] ([column1], [column2])

Unique Clustered Index



- We know **Key values should be unique**
- How RDBMS resolves this problem?

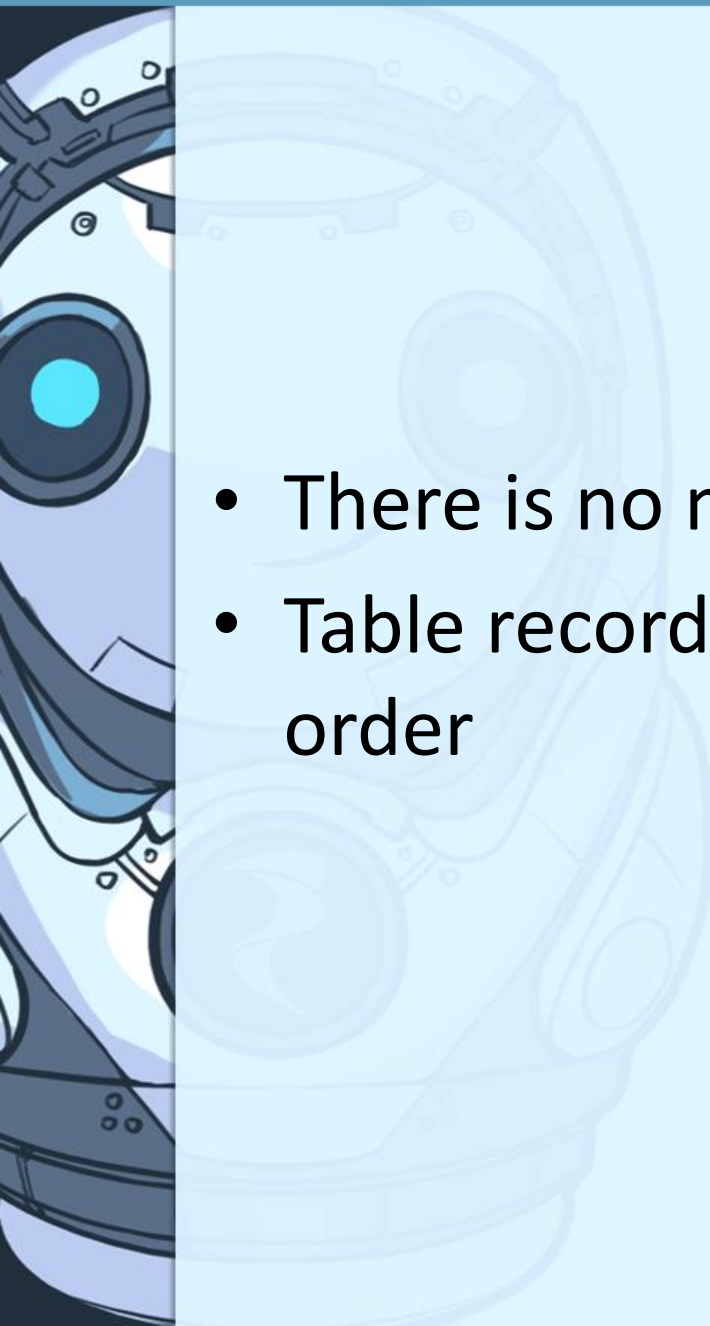
- SQL Server adds 4-byte uniquifier to each duplicated key value
- Algorithms could differ from vendor to vendor
- But the principle is the same – add something to make them unique

- Just omitting Unique keyword makes Key values **bigger** (why it's bad realize later)
- The simple truth is that Each table should have **Clustered Index**
- The Clustered Index should be always **Unique**
- The situations when its not so should be exceptional

- Such tables are called **Heap Tables**
- How are they stored in database if they do not have a Key value specified?

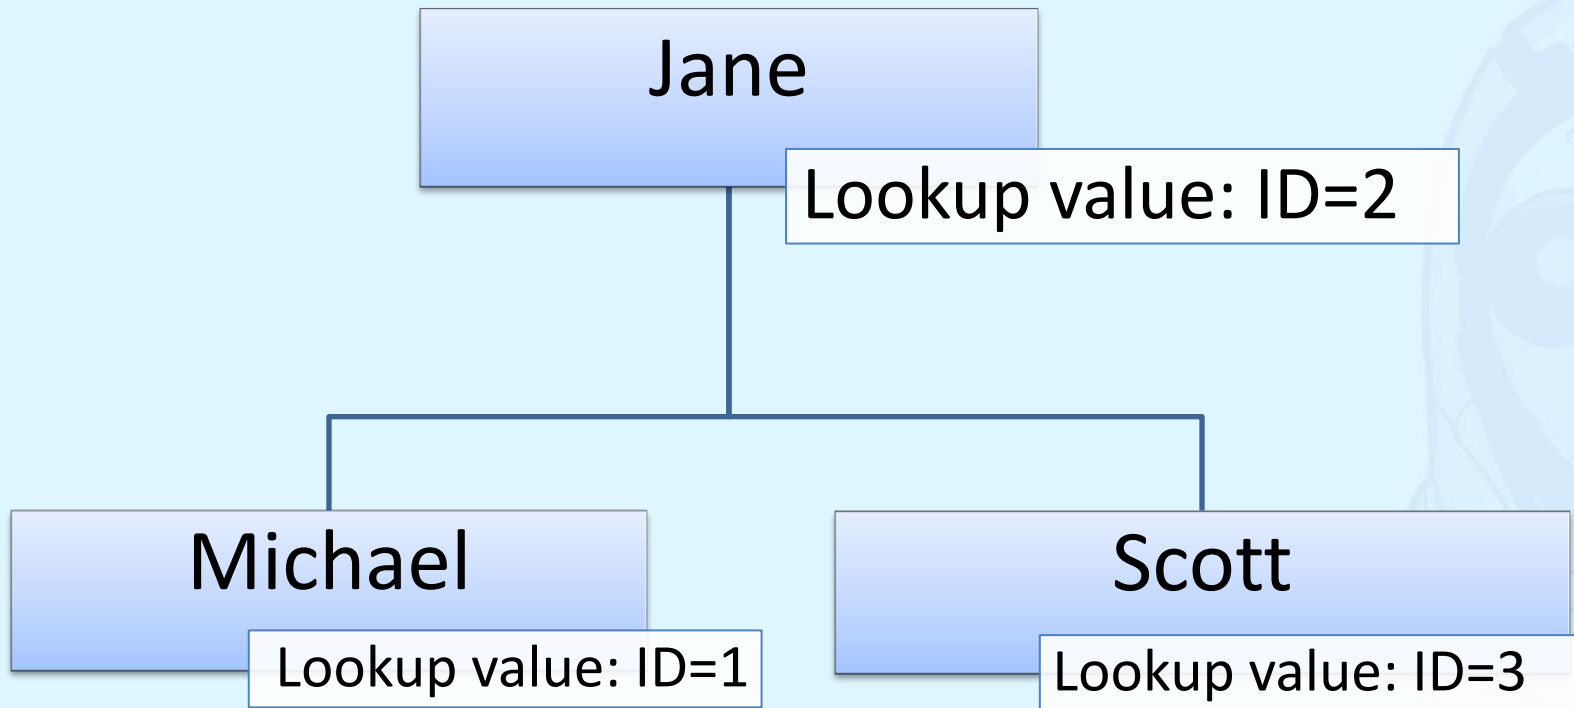
- Heap Tables are also stored in Trees
- What's in a Key value for Tables without Clustered Index?
- The value called RID
- the unique identifier which refers to the physical location of the record in a file

Why Heap Tables are so bad

- 
- There is no meaningful data in Keys
 - Table records are not stored physically in Keys' order


- Clustered Index has the actual data columns in Leaf-nodes
- What's in Leaf-node of Non-clustered index?
- Remember that **Non-clustered Indexes are duplicated data**

- Leaf-nodes contain the lookup values
- Lookup value is **Clustered Index's Key**



- We know **Key values should be unique**
- How non-clustered index's key becomes unique?

- SQL Server adds **Clustered Index Key value** to Non-clustered Index Key value to make it unique



- from SELECT statement the WHERE condition is taken
- based on the Columns in WHERE we know what columns we search by
- look through available indexes trying to find the appropriate one, starting from Clustered
- found out non-clustered index which fits best

- get the needed Node in Non-clustered index
- get the Lookup value from that Node
- use that lookup value to find a record in Clustered index
- get selected columns from Clustered index (table itself)

- Unique Clustered Index on **Id** column
- Non-unique Non-clustered Index on **City** column
- Select UserName from tbl where City = 'Boston'

Id	User Name	Country	City	Age
1	Michael	USA	Boston	30
2	Jane	USA	Boston	24
3	Scott	USA	NYC	18
4	Bob	UK	London	41
5	Prescott	UK	London	35

- Unique Clustered Index on **Id** column
- Non-unique Non-clustered Index on **City** column
- Select **Id** from tbl where City = 'Boston'

Id	User Name	Country	City	Age
1	Michael	USA	Boston	30
2	Jane	USA	Boston	24
3	Scott	USA	NYC	18
4	Bob	UK	London	41
5	Prescott	UK	London	35

- Unique Clustered Index on **Id** column
- Non-unique Non-clustered Index on **City** column
- Select **UserName** from tbl where City = 'Boston'
select should not go to Clustered Index

Id	User Name	Country	City	Age
1	Michael	USA	Boston	30
2	Jane	USA	Boston	24
3	Scott	USA	NYC	18
4	Bob	UK	London	41
5	Prescott	UK	London	35

- Unique Clustered Index on **Id, UserName** column
- Select **Id** from tbl where City = 'Boston' and UserName = 'Michael'
- What columns Non-unique Non-clustered Index would include?

Id	User Name	Country	City	Age
1	Michael	USA	Boston	30
2	Jane	USA	Boston	24
3	Scott	USA	NYC	18
4	Bob	UK	London	41
5	Prescott	UK	London	35

***Follow us on Twitter
@Sperasoft***

***Visit our site:
sperasoft.com***