







# SQL Server Service Broker Scale out by Reliable Messaging

- 
- Scalability – what does Scale Out means?
  - Sharding
  - Problems of sharding
  - How does Service Broker help?


- 
- Is the ability of a system to **handle growing amount** of work in a capable manner
  - or its **ability to be enlarged** to accommodate that growth

<http://en.wikipedia.org/wiki/Scalability>

- 
- **Scale up** – scale vertically
  - **Scale out** – scale horizontally

- 
- means **vertically**
  - add resources to a single node in a system
    - CPU
    - memory
    - storage
    - etc



- 
- means **horizontally**
  - add more nodes to a system
  - in databases that's called **Sharding**
  - Microsoft used another term in SQL Azure - Federations

- Twitter Gizzard -  
<https://github.com/twitter/gizzard>
- written in Scala



**Sharding** involves two technologies:

- Partitioning
- Replication



- Data is divided into chunks
- Stored across computers
- Each chunk is small enough that we can **manipulate and query efficiently**



User = “Denis”



User = “Eva”

- multiple copies of data are stored across several machines
- efficiently responds to tons of queries
- resilient to failures



User = "Eva"  
User = "Denis"





User = "Eva"  
User = "Denis"

Besides Scaling Out:

- Disaster Recovery strategies have You to make database backups
- Replicated shards are working copies
- Need a new backup? Replicate a new Shard!

# Are there any Problems?

- 
- Yes! **Sharding is difficult!**
  - We have to think of how to partition our data correctly
  - How to ensure that all copies of data are consistent?
  - Write conflicts – no guarantee that operations will apply in order

- 
- “Shared Nothing Architecture”
  - no ongoing need to retain shared access between shards
  - each Shard can live in a totally separate
    - instance
    - database server
    - data center
    - continent



# Partitioning Problem

- Orders table is partitioned by Region
- Orders references to Users
- Users is **not partitioned**, it is **replicated**
- How to maintain Users in a consistent state?



Orders [Region = “**US**”]  
Users (not partitioned)



Orders [Region = “**Russia**”]  
Users (not partitioned)

Front-End

Denis writes a new Status

POST



User [Name = "Denis"]  
Wall Statuses

How do we get Denis' statuses?

We should see that status  
on Eva's wall

GET



User [Name = "Eva"]  
Wall Statuses

# Possible Solution 1

Front-End

Denis writes a new Status

POST



User [Name = "Denis"]  
Wall Statuses

**No benefit from Shards then! Wrong!**

Duplicate the status?

POST



User [Name = "Eva"]  
Wall Statuses

# Possible Solution 2

Front-End

Denis writes a new Status

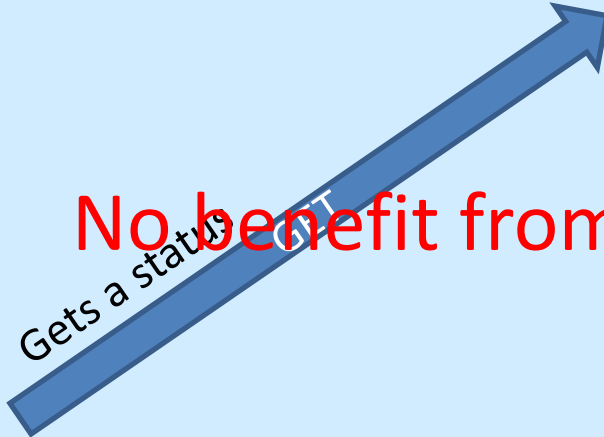
POST



User [Name = "Denis"]  
Wall Statuses

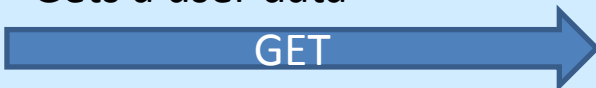
Gets a status

GET



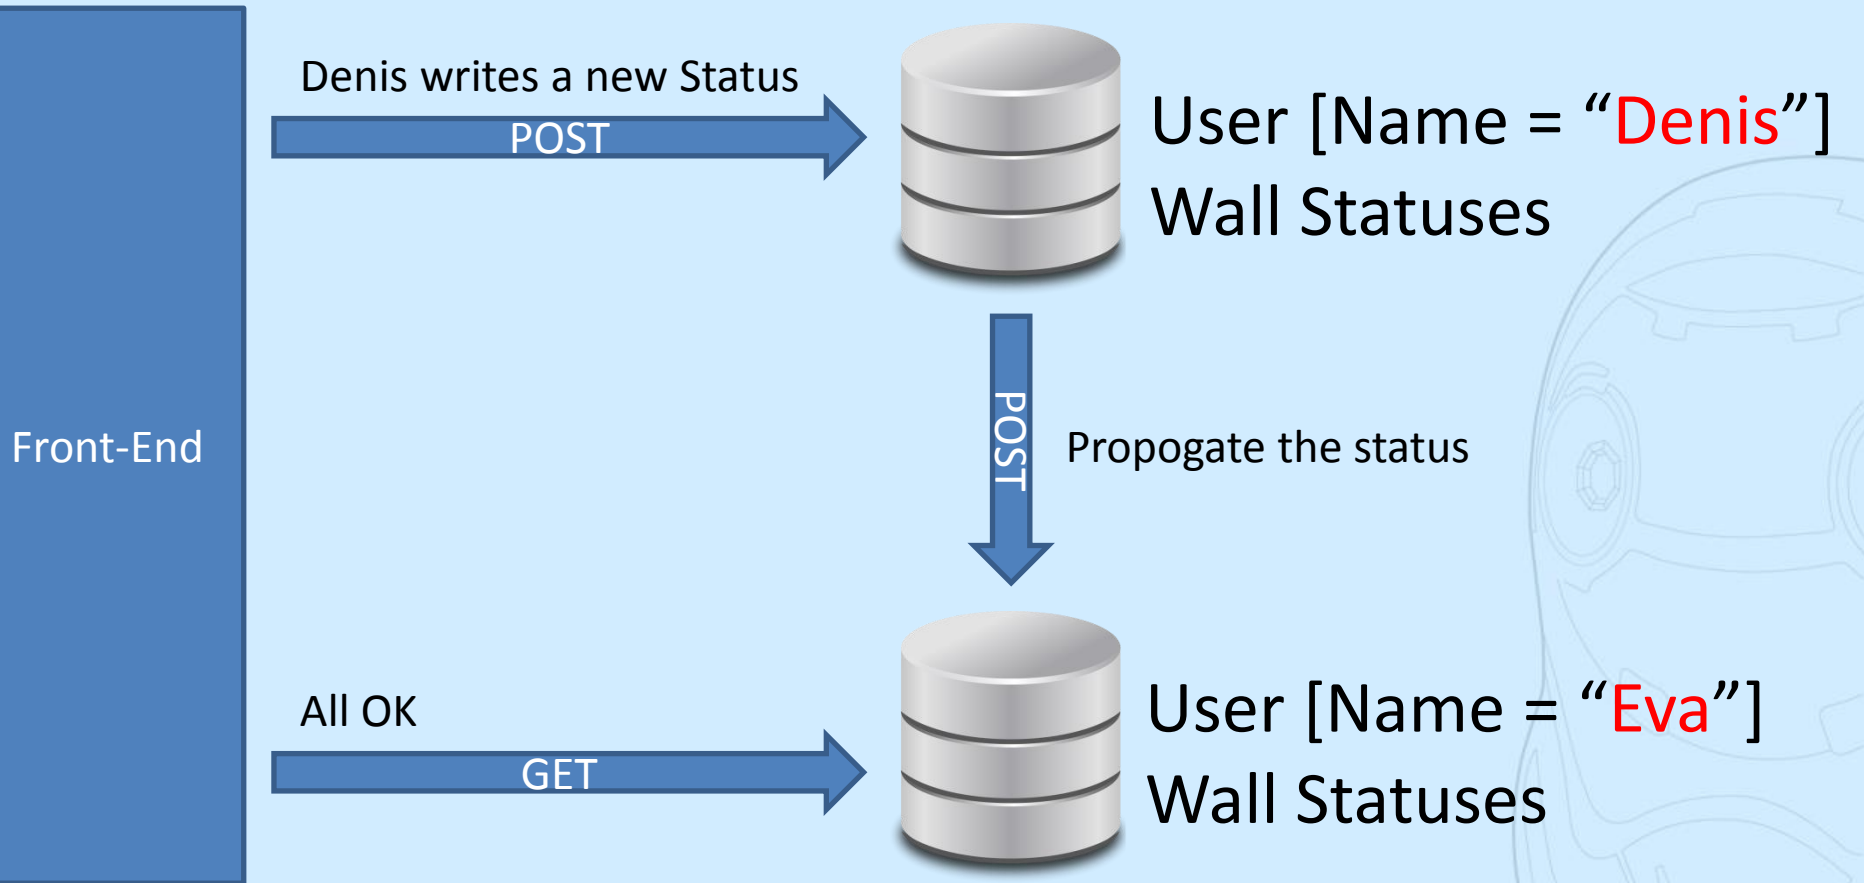
Gets a user data

GET



User [Name = "Eva"]  
Wall Statuses

**No benefit from Shards then! Wrong!**





- custom application which reads from one Shard and writes to the others
- add all Shards as Linked servers, write stored procedures to write a new record to each remote server
- SQL Server Service Broker

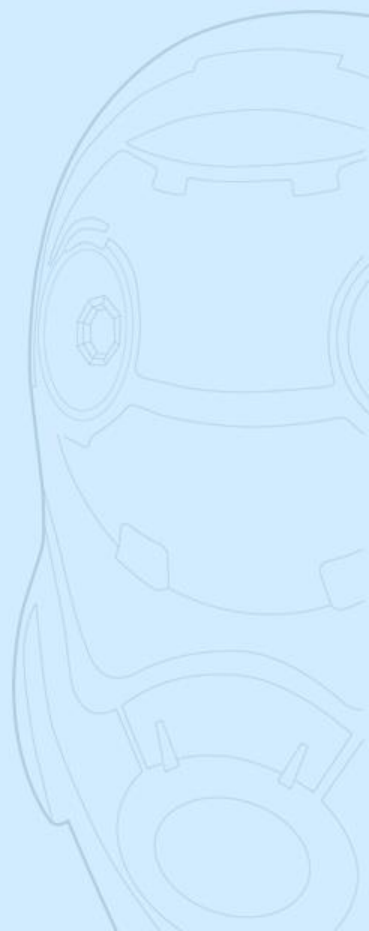
- allows **Internal** and **External** processes to send and receive guaranteed asynchronous messages
- messages are sent to Queues in the same database, same instance, same server, another database, another instance, remote server

the largest known implementation is in **MySpace**

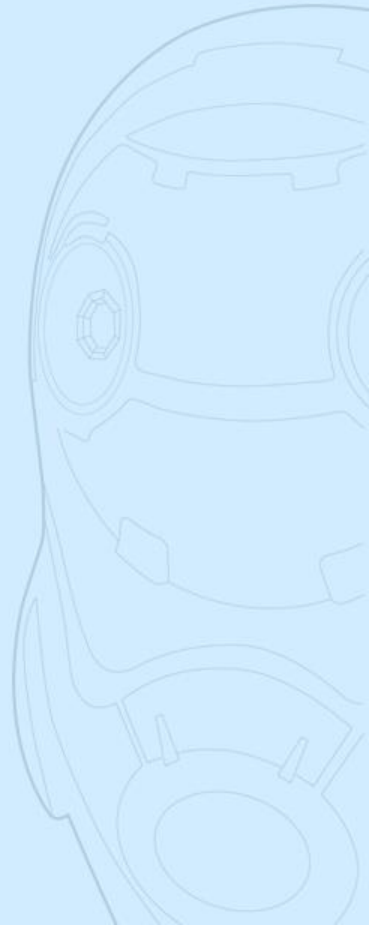
- **440** SQL Server instances
- over **1,000** databases
- **1 Petabyte** of data (1 million gigabytes)

Performance:

- **5,000** messages/second in Labs
- **18,000** messages/second in MySpace



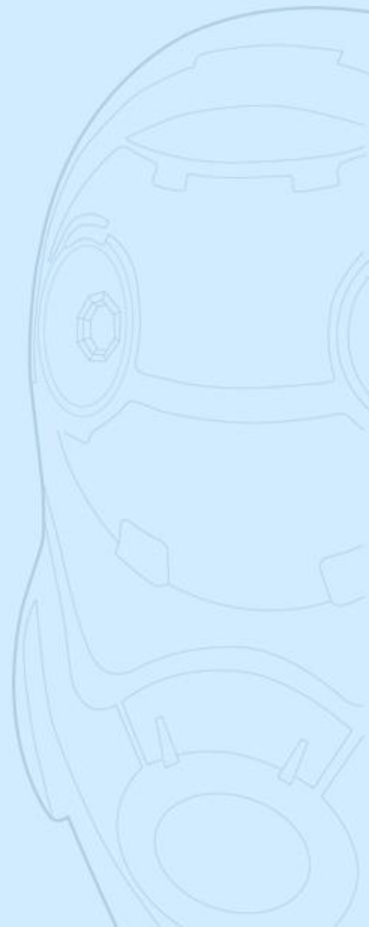
- Always Transactional
- Asynchronous
- Queued
- All messages are in XML format
- Routing
- Multicast messages in SQL Server 2012



- No Administration Tools – not a lack - no at all
- Queue is a Table in database – index fragmentation when we cannot empty queue
- complicated T-SQL syntax and object model



- Message Types
- Contracts
- Queues
- Services
- Routes
- Activation
- Remote Service Bindings
- Dialog Conversation
- Conversation Groups



- Dialog Security – between two Services
  - encrypts messages in a dialog conversation
  - provides the remote authorization
- Transport Security
  - establishes an authenticated network connection between two Databases

- Enable Service Broker at Database level

```
ALTER DATABASE SampleDatabase  
    SET ENABLE_BROKER;  
GO
```

- Always enable Service Broker in **MSDB** system database in each SQL Server instance
- Enable SQL Server to communicate to another instance at Database level

```
ALTER DATABASE SampleDatabase  
    SET TRUSTWORTHY ON;  
GO
```

- *Create Send Message type*

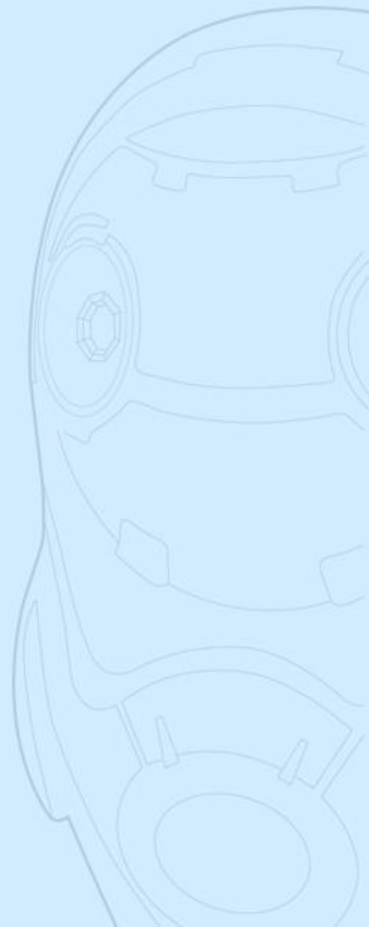
```
CREATE MESSAGE TYPE NewUserRequest  
VALIDATION = WELL_FORMED_XML;
```

- *Create Receive Message Type*

```
CREATE MESSAGE TYPE NewUserResponse  
VALIDATION = WELL_FORMED_XML;
```

- create Contract between two services

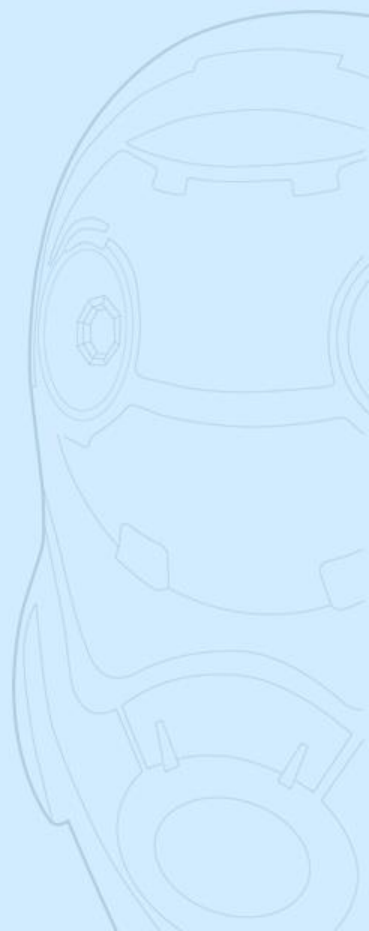
```
CREATE CONTRACT NewUserContract (  
    NewUserRequest SENT BY INITIATOR,  
    NewUserResponse SENT BY TARGET  
);
```





```
CREATE QUEUE NewUserReceiveQueue  
  WITH STATUS = ON,  
  ACTIVATION (  
    PROCEDURE_NAME = OnReceiveNewUser,  
    MAX_QUEUE_READERS = 5,  
    Execute AS 'dbuser'  
  );
```

```
CREATE QUEUE NewUserSendQueue  
  WITH STATUS = ON;
```



- Service will be listening for messages in a queue and react only on those which apply to the specified contract

```
CREATE SERVICE NewUserSendService  
  ON QUEUE NewUserSendQueue (  
    [NewUserContract]  
  );
```

```
CREATE SERVICE NewUserReceiveService  
  ON QUEUE NewUserReceiveQueue (  
    [NewUserContract]  
  );
```



Sender:

- begins a transaction
- begins a dialog conversation from **NewUserSendService** to **NewUserReceiveService**
- sends one or more XML messages to that dialog
- commits a transaction
- Waits for a Reply message
- Processes Reply messages when arrived

## Receiver (Target):

- when a new message is delivered to Receiver's **NewUserReceiveQueue** the **OnReceiveNewUser** stored procedure is activated by Service Broker
- Stored procedure begins a transaction
- receives messages out of the **NewUserReceiveQueue** queue
- sends Reply messages
- ends a dialog conversation
- Do what you want with that message
- commits a transaction